


Modeled Decode Latency — TieredKV vs. SOTA (Fair / Iso-VRAM)



Latency: analytical model (GPU VRAM $\leftrightarrow$ STT-RAM tier bandwidth). TieredKV latency includes migration overhead.