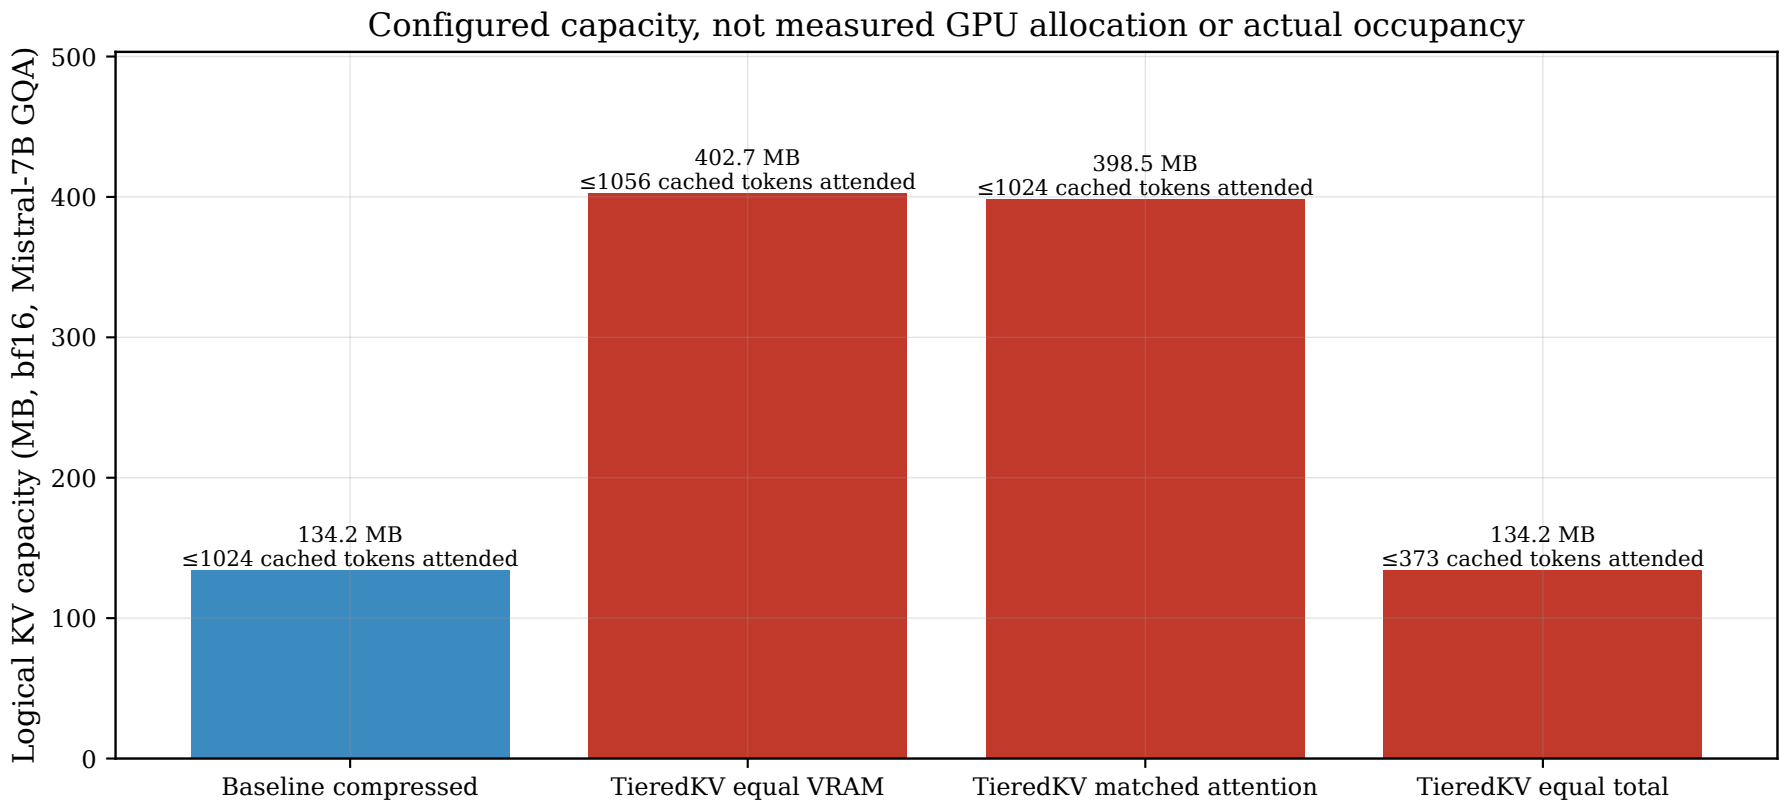
Configured capacity, not measured GPU allocation or actual occupancy
Logical KV capacity (MB, bf16, Mistral-7B GQA)
500
400
300
200
100
0
134.2 MB
≤1024 cached tokens attended
402.7 MB
≤1056 cached tokens attended
398.5 MB
≤1024 cached tokens attended
134.2 MB
≤373 cached tokens attended
Baseline compressed
TieredKV equal VRAM
TieredKV matched attention
TieredKV equal total