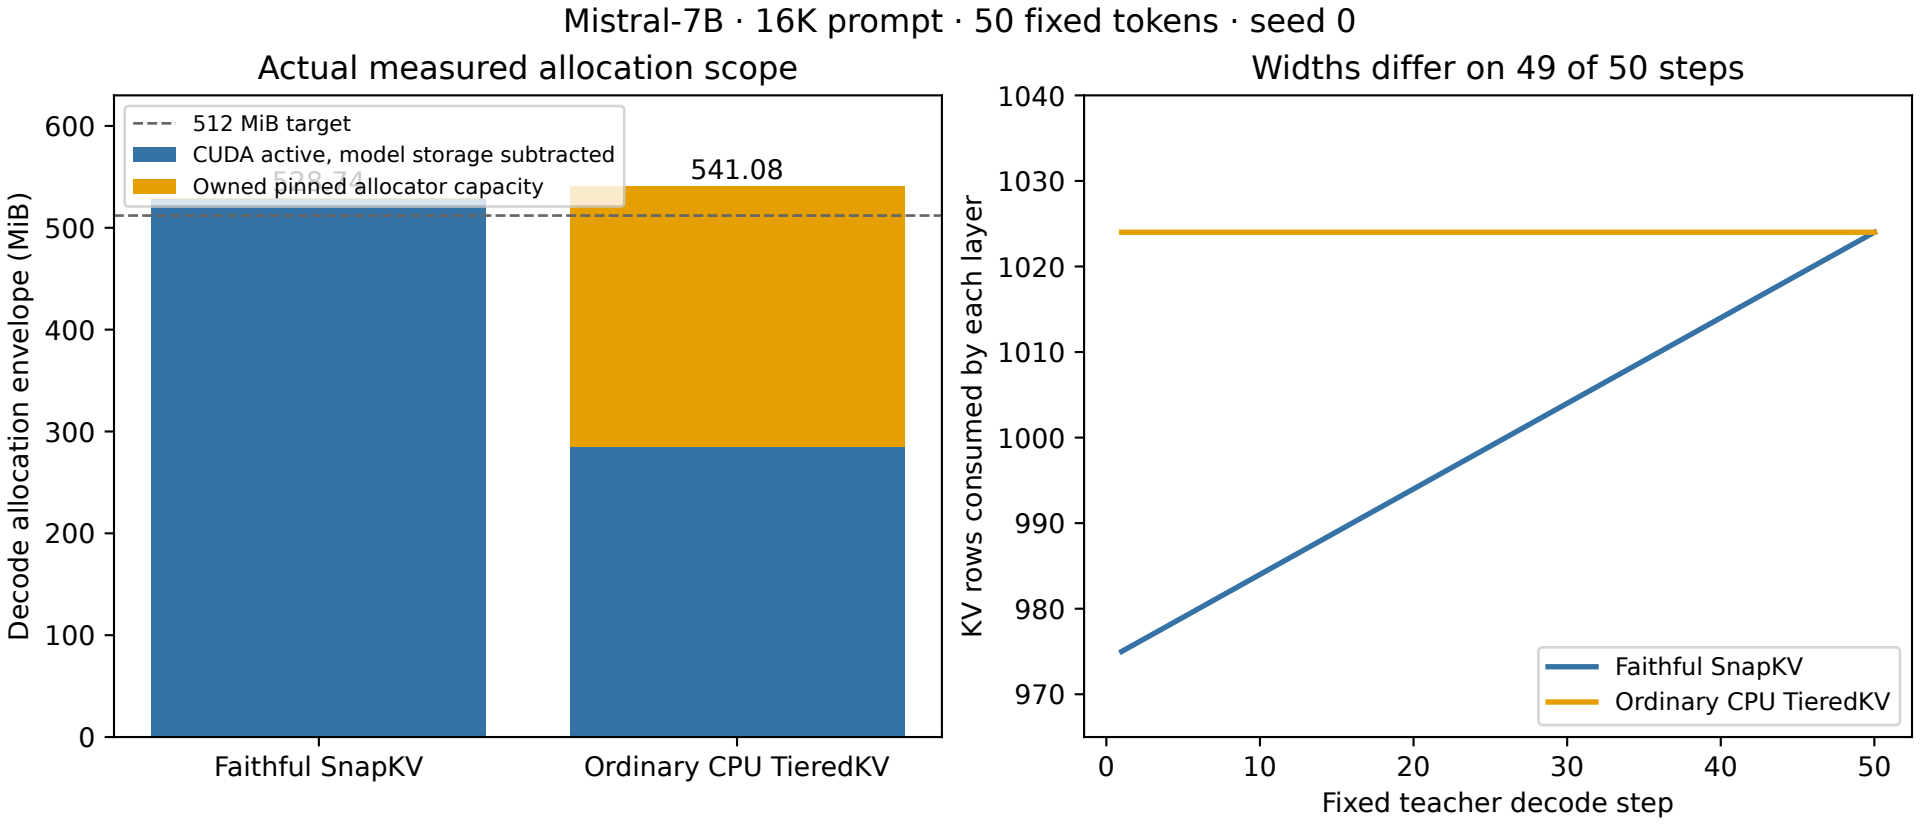
Mistral-7B · 16K prompt · 50 fixed tokens · seed 0
Actual measured allocation scope
512 MiB target
CUDA active, model storage subtracted
Owned pinned allocator capacity
541.08
Decode allocation envelope (MiB)
0
100
200
300
400
500
600
Faithful SnapKV
Ordinary CPU TieredKV
Widths differ on 49 of 50 steps
KV rows consumed by each layer
970
980
990
1000
1010
1020
1030
1040
0
10
20
30
40
50
Fixed teacher decode step
Faithful SnapKV
Ordinary CPU TieredKV